

# APACHE AIRFLOW

- Apache Airflow is an open-source platform used to programmatically author, schedule, and monitor workflows. It is widely used for orchestrating complex computational workflows and data pipelines.
- **Key Features:**
  - Workflow orchestration using Directed Acyclic Graphs (DAGs)
  - Dynamic pipeline generation using Python
  - Scalability and extensibility
  - Integration with various databases, cloud services, and data tools
  - Real-time monitoring and logging of tasks

## Architecture

Apache Airflow follows a modular architecture. Its key components include:

- **Scheduler:**

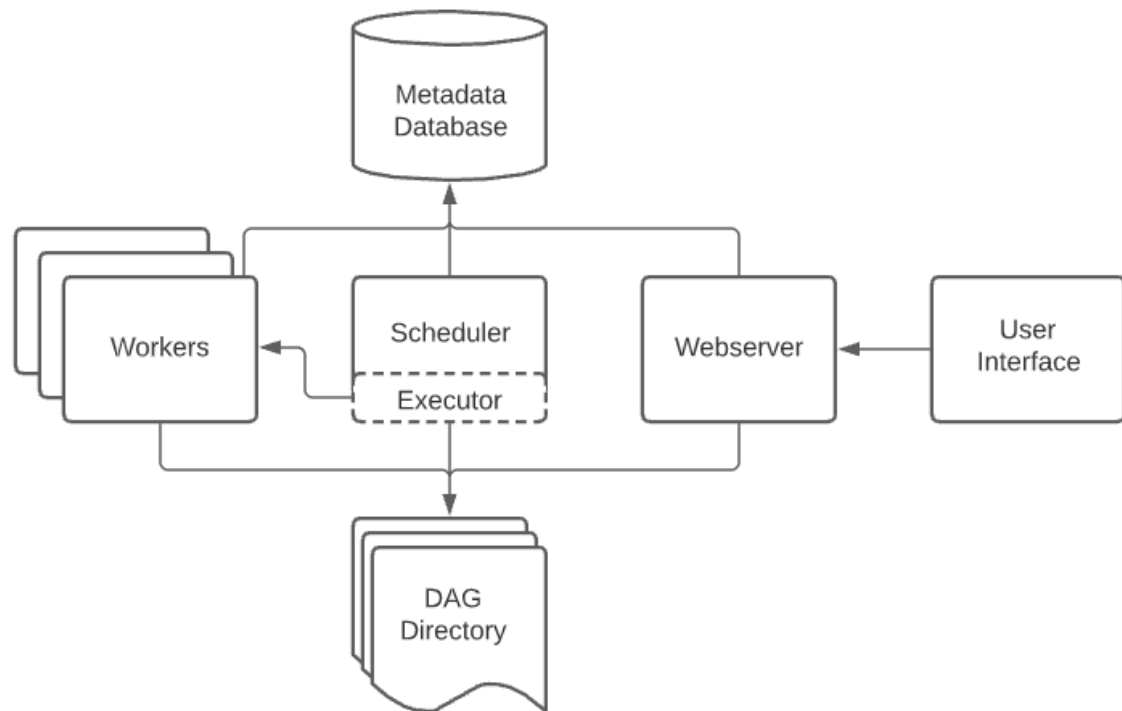
Responsible for scheduling DAGs and triggering tasks based on dependencies and timing.
- **Executor:**

Executes the tasks. Types include SequentialExecutor, LocalExecutor, CeleryExecutor, and KubernetesExecutor.
- **Metadata Database:**

Stores DAG definitions, task status, and runtime information. Commonly uses PostgreSQL or MySQL.
- **Web Server:**

Provides a UI to monitor workflows, view logs, and manage DAGs.
- **Workers:**

Execute tasks in distributed setups (used with Celery/Kubernetes executors).



## CONCEPTS :

### 1. DAG (Directed Acyclic Graph):

- Represents the workflow. Each node is a task, and edges represent dependencies.

### 2. Task:

- A single unit of work in a DAG. Example: running a Python script, transferring data, or executing SQL queries.

### 3. Operator:

- Predefined task templates. Types include:
  - PythonOperator: Execute Python functions
  - BashOperator: Run bash commands
  - EmailOperator: Send emails
  - PostgresOperator: Execute SQL on PostgreSQL

### 4. Task Instance:

- A specific run of a task at a particular time.

### 5. Scheduler Interval:


- Defines how often a DAG should run (e.g., @daily, @hourly).

### 6. Hooks and Connections:

- Connectors to databases, cloud services, and APIs.

## DAGs

- A *DAG* (Directed Acyclic Graph) is the core concept of Airflow, collecting [Tasks](#) together, organized with dependencies and relationships to say how they should run.



- It defines four Tasks - A, B, C, and D - and dictates the order in which they have to run, and which tasks depend on what others.
- It will also say how often to run the DAG - maybe “every 5 minutes starting tomorrow”, or “every day

## Use Cases

- ETL (Extract, Transform, Load) pipelines
- Data warehouse loading
- Machine learning model training pipelines
- Batch job scheduling
- Data quality monitoring

## Advantages:

- Flexible and Python-based
- Scalable and extensible
- Active community support
- Rich UI for monitoring

## Limitations:

- Learning curve for beginners
- Complex setup in distributed environments
- Not suitable for real-time streaming tasks (use Apache Kafka or Spark Streaming for real-time)